





Accuracy vs. visual-token budget, attention top-K against unranked floors

LLaVA-OneVision-7B



Qwen2.5-VL-7B



- attn top-K (early)
- attn top-K (mid)
- attn + MMR diversity (mid)
- random
- uniform
- uniform (staggered)
- attn bottom-K (early)
- attn bottom-K (mid)
- no pruning (per panel)
- text only, 0 visual tokens (per panel)